

# An Introduction to Machine Learning and Deep Learning with R



HARPO (AKA Carlos A. Catania Ph.D.)  
LABSIN - Ingeniería - UNCuyo



@harpolabs



harpo@ingenieria.uncuyo.edu.ar



## AGENDA Day 2: Deep Learning



- Basic concepts
- Neural networks
- Convolutional Neural Networks
- Recurrent Neural networks (??)



## Materials Day 2: Deep Learning

- **Rstudio** desktop or server version  $\geq 1.1$
- **R** version  $\geq 3.5$
- **Rstudio** Notebook operational
- `install.packages("keras")`
- `install_keras()`
- Data for day 2 at <http://bit.ly/mlab2019>

# Highly recommended books

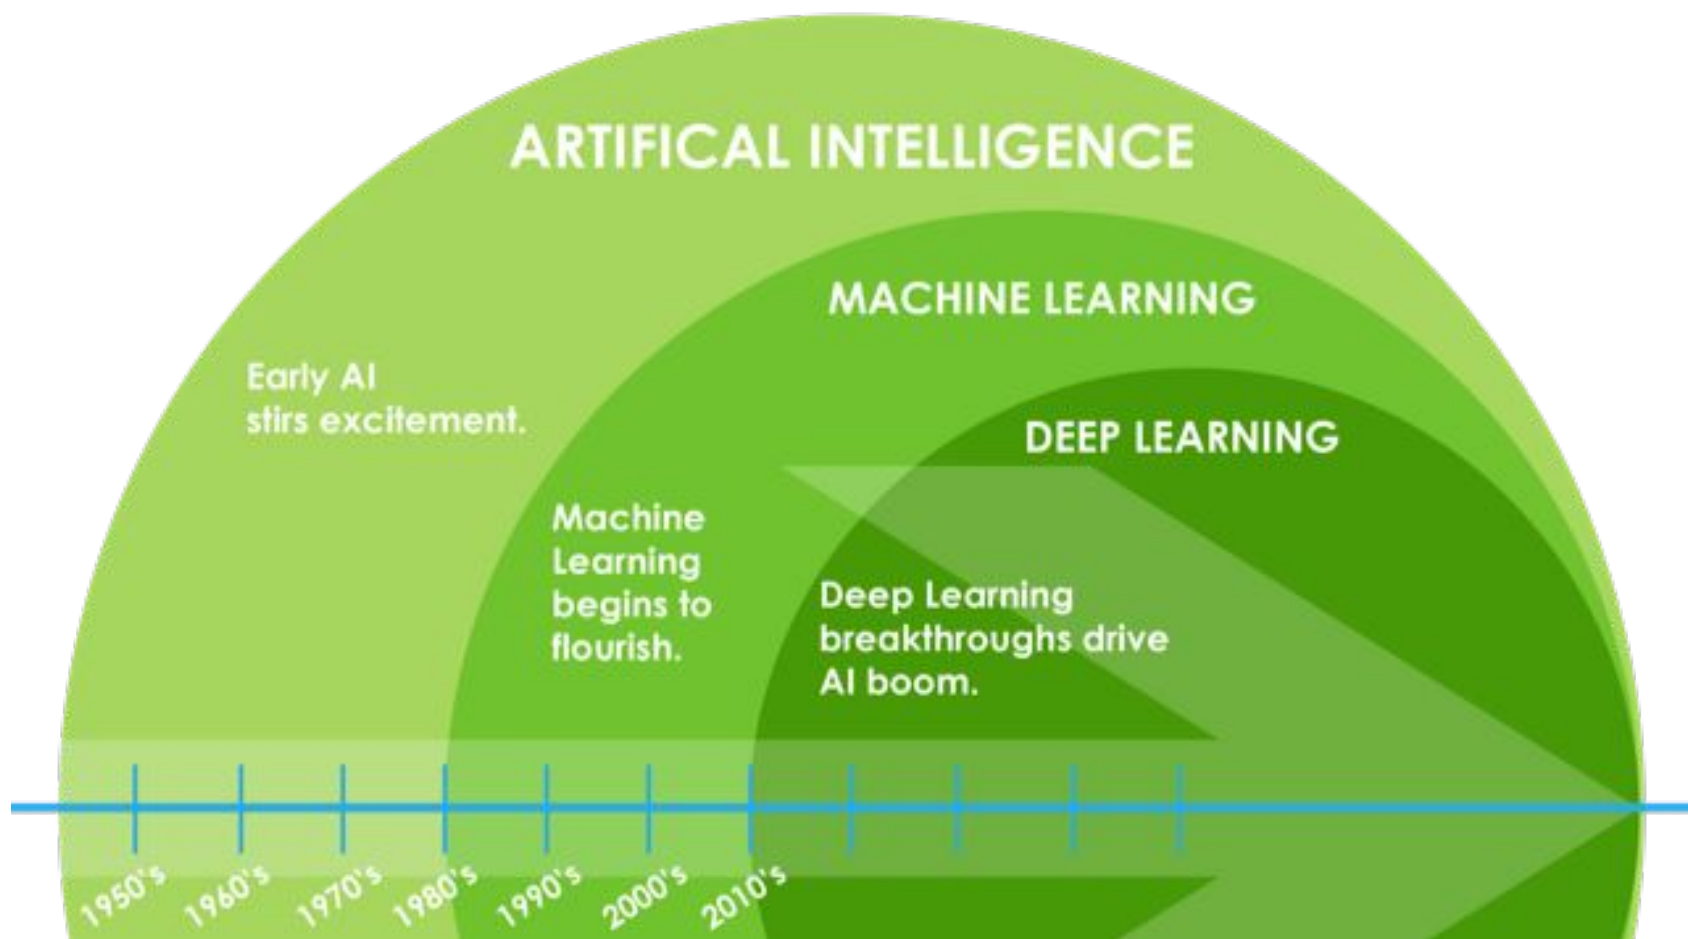




# Deep Learning



# Deep Learning (in 10 min??)



---

---

# DEEP LEARNING

What is **LEARNING**?


What is **DEEP**?

---




---

# The general idea behind ML



---

# A simple example



**Goal:** predict data point color based on its coordinate  $(x,y)$


**Method:** We need a new representation of data that cleanly separates the white points from the black points.

---


# Simple approach: Coordinates Change

The problem can be expressed as a simple rule: “**Black points** are such that  $x > 0$ ,” or “**White points** are such that  $x < 0$ .”


1: Raw data



2: Coordinate change



3: Better representation



---

# LEARNING


- In the context of **machine learning**, describes an automatic search process for better representations.
- In a way, all machine-learning algorithms consist of **automatically** finding such **transformations** that turn data into more-useful representations for a given task.

---

# However...

Sometimes we need to **feed** the algorithm with some particular presentation aiming at helping the final performance.

Such process is usually call **feature engineering** and could be **extremely time consuming**



---


# The “deep” in Deep Learning

Deep learning is a specific subfield of machine learning: a new take on learning representations from data that puts an emphasis on learning successive *layers* of increasingly meaningful representations.


---

# Neural networks

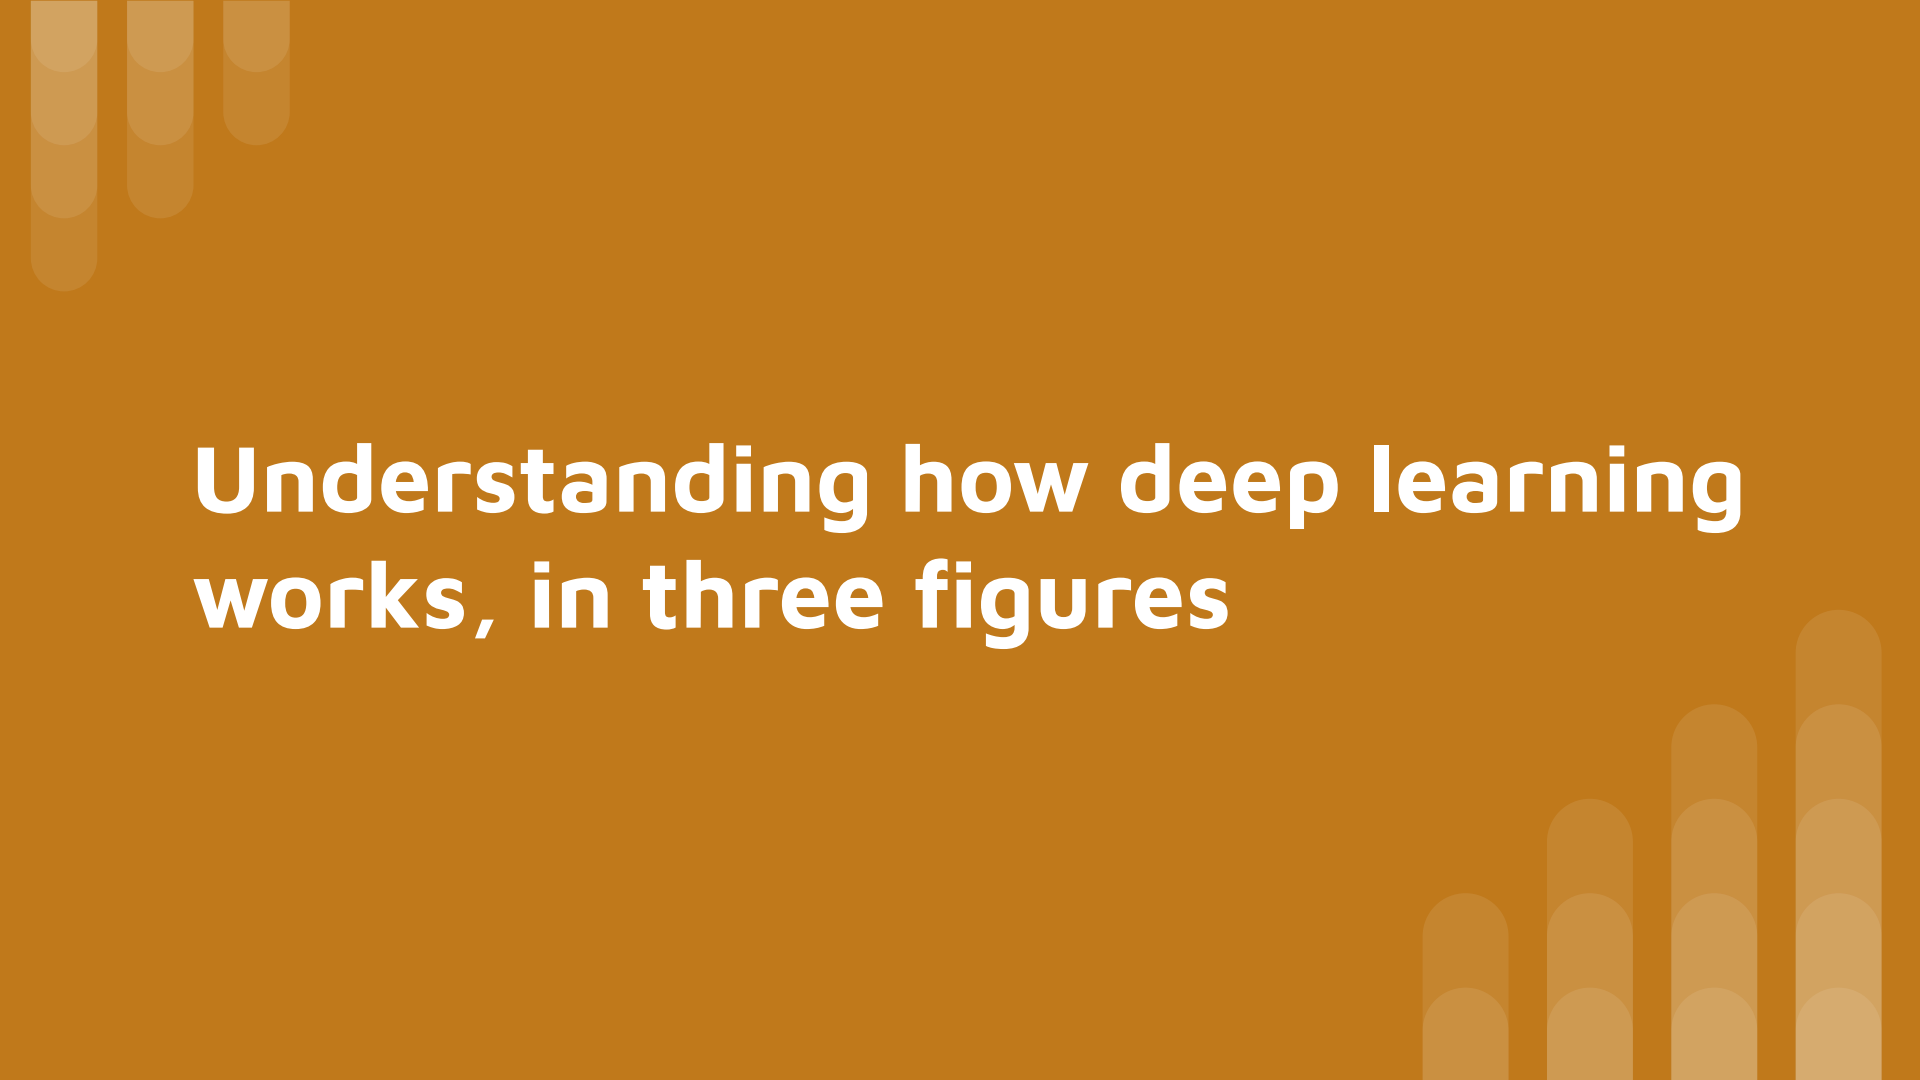
In deep learning, these layered representations are (almost always) learned via models called *neural networks*



# Deep learning for digit classification



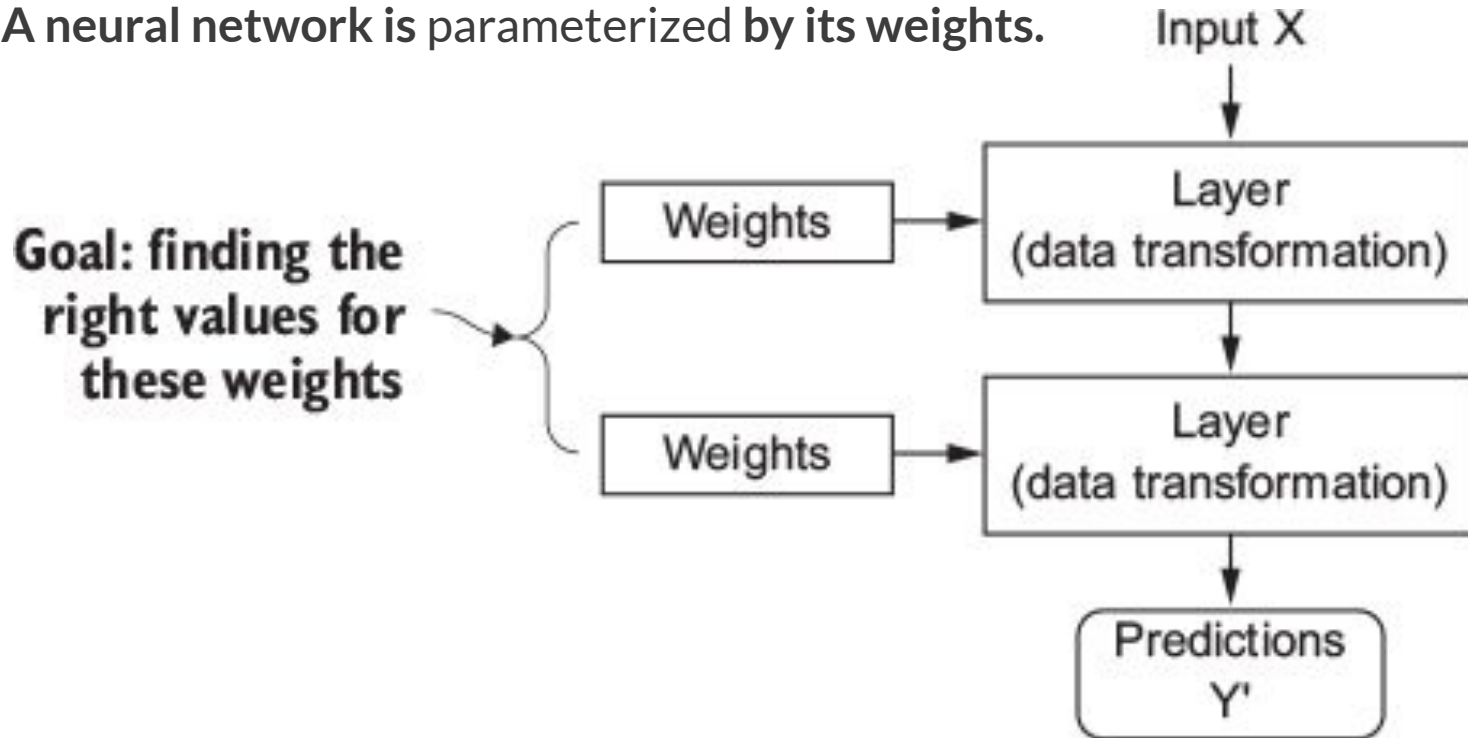




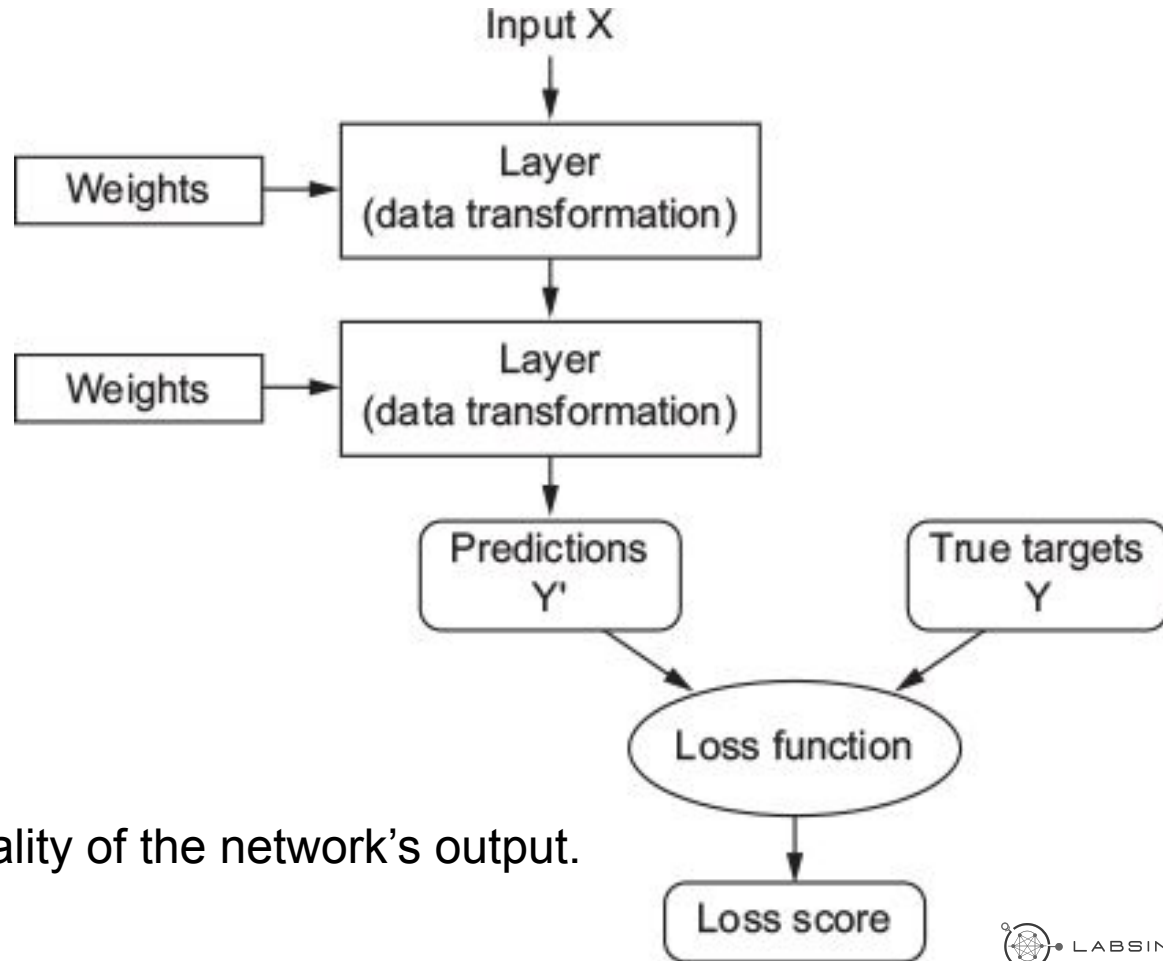
# Understanding how deep learning works, in three figures

# Figure 1:

A neural network is parameterized by its weights.



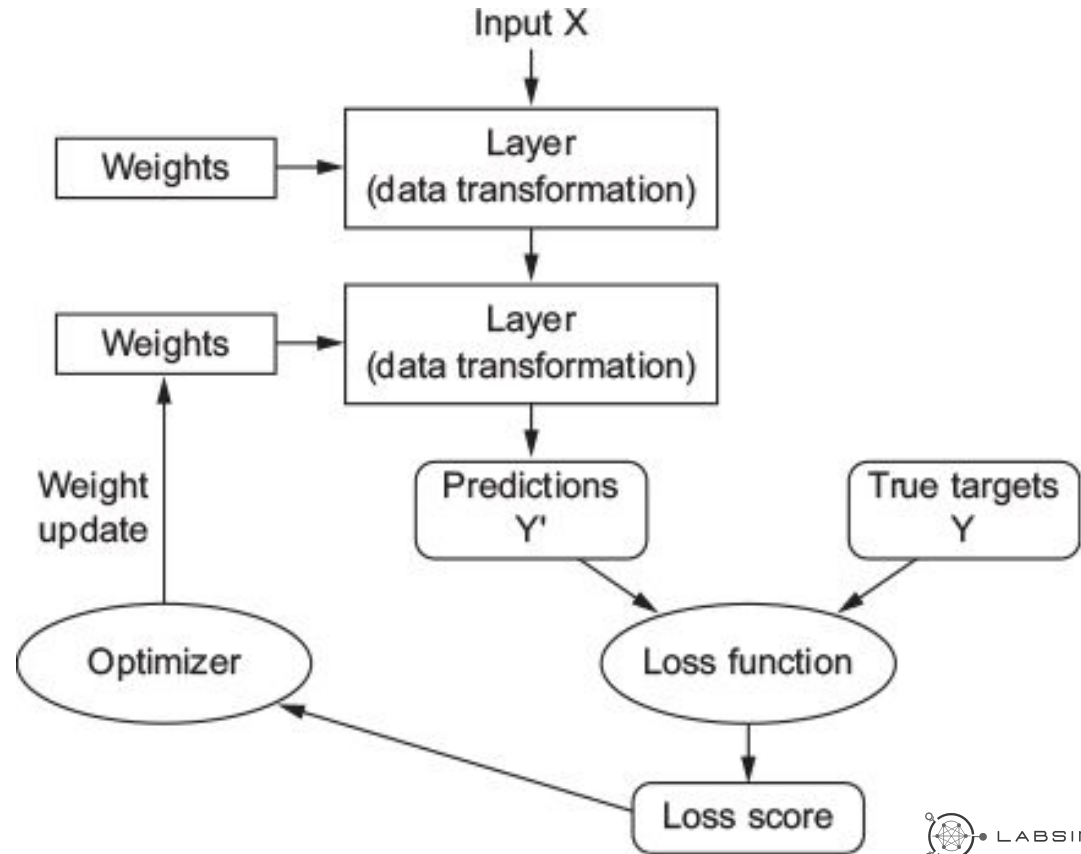
# Figure 2:



A loss function measures the quality of the network's output.

# Figure 3:

The loss score is used as a feedback signal to adjust the weights.



# LAB 3: The MNIST dataset

The problem we're trying to solve here is to classify grayscale images of handwritten digits (28 pixels by 28 pixels) into their 10 categories (0 to 9).




# Loading the Mnist Dataset

```
```{r, results='hide'}
library(keras)

mnist <- dataset_mnist()
train_images <- mnist$train$x
train_labels <- mnist$train$y
test_images <- mnist$test$x
test_labels <- mnist$test$y
```
```

```
```{r}
str(train_images)
```
```

```
int [1:60000, 1:28, 1:28] 0 0 0 0 0 0 0 0 0 0
```



The images are encoded as 3D arrays, and the labels are a 1D array of digits, ranging from 0 to 9.

# Let's build our first model

```
```{r}
network <- keras_model_sequential() %>%
  layer_dense(units = 512, activation = "relu", input_shape = c(28 * 28)) %>%
  layer_dense(units = 10, activation = "softmax")
```
```

- The core building block of neural networks is the **\_layer\_**, a data-processing module that you can think of as a filter for data. Some data comes in, and it comes out in a more useful form.
- Specifically, layers extract **\_representations\_** out of the data fed into them—hopefully representations that are more meaningful for the problem at hand.

# Neural Networks Fully Connected

Color Guided Matrix Multiplication for a Binary Classification Task  
with  $N = 4$

**Input Layer**

bias X1 X2


$$\begin{bmatrix} 1 & 0 & 1 \\ 1 & 0 & 0 \\ 1 & 0 & 0 \\ 1 & 1 & 0 \end{bmatrix} \cdot \begin{bmatrix} .5 & .5 & .5 \\ .5 & .5 & .5 \\ .5 & .5 & .5 \end{bmatrix} =$$

$4 \times 3$   $3 \times 3$

Weights  $w^T$  (transposed)

Go to Hidden Nodes

1 2 3



**Hidden Layer**

Bias

Node 1

Node 2

Node 3

$$= \begin{bmatrix} 1 & 1 & 1 \\ .5 & .5 & .5 \\ .5 & .5 & .5 \\ 1 & 1 & 1 \end{bmatrix} \cdot \begin{bmatrix} .5 & .5 & .5 \\ .5 & .5 & .5 \\ .5 & .5 & .5 \end{bmatrix} =$$

$4 \times 3$

**Sigmoid Function**

$\frac{1}{1 + e^{-(wx+b)}}$

**Weights**

$3 \times 2$

**Output Layer**

$4 \times 2$

**Sigmoid Function**

$\frac{1}{1 + e^{-(wx+b)}}$


**Output**

$$= \begin{bmatrix} 1 & 0 \\ 1 & 0 \\ 1 & 0 \\ 1 & 0 \end{bmatrix}$$

Rubens Zimbres



# Softmax activation function



# Let's build our first model

To make the network ready for training, we need to pick three more things:

- **A loss function:** How the network will be able to measure how good a job it's doing on its training data, and thus how it will be able to steer itself in the right direction.
- **An optimizer:** The mechanism through which the network will update itself based on the data it sees and its loss function.
- **Metrics** to monitor during training and testing: i.e. accuracy (the fraction of the images that were correctly classified).

```
network %>% compile(  
  optimizer = "rmsprop",  
  loss = "categorical_crossentropy",  
  metrics = c("accuracy")  
)
```

# But before.. One more thing..

Before training, we'll preprocess the data by reshaping it into the shape the network expect


```
train_images <- array_reshape(train_images, c(60000, 28 * 28))  
train_images <- train_images / 255  
  
test_images <- array_reshape(test_images, c(10000, 28 * 28))  
test_images <- test_images / 255
```

# LAB 3: Regression, the Boston Dataset.

1. What is the output of the Keras model when using regression?
2. Callbacks
3. Scaling vs. no Scaling.


# Convolutional Neural Networks: (CNN)

Any NN including a Convolutional Layer



1. **Dense layers learn global patterns** (for example, for an MNIST digit, patterns involving all pixels)
2. **Convolution layers learn local patterns**

# The hierarchical approach of CNN



# Convolutional Neural Networks in KERAS

```
layer_conv_2d(filters = 64, kernel_size = c(3,3), activation = 'relu')
```

Convolutions operate over 3D tensors, called **feature maps**, with two spatial axes (*height* and *width*) as well as a depth axis (also called the *channels* axis).

- For an **RGB image**, the **dimension of the depth axis is 3**, because the image has three color channels: red, green, and blue.
- For a **black-and-white picture**, like the MNIST digits, the **depth is 1** (levels of gray)

# Convolutional Neural Networks

- This output feature map is still a 3D tensor: it has a width and a height.
- The output depth is a parameter of the layer, **the so called filters**.
- Filters encode specific aspects of the input data: at a high level, a single filter could encode the concept “presence of a face in the input,” for instance.



# The Convolution Operation

|                 |                 |                 |   |   |
|-----------------|-----------------|-----------------|---|---|
| 1 <sub>x1</sub> | 1 <sub>x0</sub> | 1 <sub>x1</sub> | 0 | 0 |
| 0 <sub>x0</sub> | 1 <sub>x1</sub> | 1 <sub>x0</sub> | 1 | 0 |
| 0 <sub>x1</sub> | 0 <sub>x0</sub> | 1 <sub>x1</sub> | 1 | 1 |
| 0               | 0               | 1               | 1 | 0 |
| 0               | 1               | 1               | 0 | 0 |

Image


|   |  |  |
|---|--|--|
| 4 |  |  |
|   |  |  |
|   |  |  |

Convolved  
Feature

# The (Max) Pooling Operation

```
layer_max_pooling_2d(pool_size = c(2, 2))
```

- Used in order to downsample the feature maps.
- Max pooling consists of extracting windows from the input feature maps and outputting the max value of each channel.
- Similar to convolution, except they're transformed via a hardcoded max/average/min tensor operation instead of kernel.



# LAB 4: The MNIST DATASET with CNN

1. Compare the performance with Dense Layers
2. The input format has changed (check NN)
3. The Flatten() “Layer”
4. Adding Layers

# LAB5: The IMDB Dataset with RNN

- Dataset of 25,000 movies reviews from IMDB, labeled by sentiment (positive/negative).
- Reviews have been preprocessed, and each review is encoded as a sequence of word indexes (integers).
- For convenience, words are indexed by overall frequency in the dataset, so that for instance the integer "3" encodes the 3rd most frequent word in the data.
- This allows for quick filtering operations such as: "only consider the top 10,000 most common words, but eliminate the top 20 most common words".

**More examples at**

<https://keras.rstudio.com/articles/examples/index.html>



Carlos A. Catania (PhD)  
(AKA **Harpo**)

LABSIN - Ingeniería - UNCuyo



@harpolabs



harpo@ingenieria.uncuyo.edu.ar



<http://labsin.org>